


Name: Sameer Almadani

Student ID: 2230001818

N-Grams

It's a probabilistic model that's trained on a corpus of text. Such a model is useful in many NLP applications including speech recognition, machine translation and predictive text input. An N-gram model is built by counting how often word sequences occur in corpus text and then estimating the probabilities.

Types of N-Grams



Use-cases of N-Grams

- Auto completion of sentences
- Auto spell check
- Voice-based personal assistant bots

Unigram

Unigram is an n-gram consisting of a single item of a sequence.

- It is commonly used to calculate the probability of finding a word in an article.

$$P(A) = \frac{\text{count of } A}{\text{count of all words}}$$

```
%pip install nltk pandas certifi emoji

import os
import certifi
import nltk

os.environ.setdefault('SSL_CERT_FILE', certifi.where())
nltk.download('punkt')
nltk.download('punkt_tab')
from nltk.util import ngrams
from nltk.tokenize import word_tokenize

string = "moses supposes his toeses are roses but moses supposes erroneously"

unigrams = ngrams(word_tokenize(string), 1)

for item in unigrams:
    print(item)
```

```

Requirement already satisfied: nltk in /Users/almadani/Desktop/nlp
lab/.venv/lib/python3.13/site-packages (3.10.3)
Requirement already satisfied: pandas in
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (3.0.0)
Requirement already satisfied: certifi in
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages
(2026.5.20)
Requirement already satisfied: emoji in /Users/almadani/Desktop/nlp
lab/.venv/lib/python3.13/site-packages (2.15.0)
Requirement already satisfied: defusedxml in
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from
nltk) (0.7.1)
Requirement already satisfied: click in
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from
nltk) (8.4.2)
Requirement already satisfied: joblib in
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from
nltk) (1.5.3)
Requirement already satisfied: regex<=2021.8.3 in
/Users/almadani/Library/Python/3.13/lib/python/site-packages (from nltk) (2026.7.19)
Requirement already satisfied: tqdm in
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from
nltk) (4.67.3)
Requirement already satisfied: numpy>=1.26.0 in
/Library/Frameworks/Python.framework/Versions/3.13/lib/python3.13/site-packages (from
pandas) (2.4.0)
Requirement already satisfied: python-dateutil>=2.8.2 in
/Users/almadani/Library/Python/3.13/lib/python/site-packages (from pandas)
(2.9.0.post0)
Requirement already satisfied: six>=1.5 in
/Users/almadani/Library/Python/3.13/lib/python/site-packages (from python-
dateutil>=2.8.2->pandas) (1.17.0)

```

[notice] A new release of pip is available: 25.2 -> 26.2.1

[notice] To update, run: /Users/almadani/Desktop/nlp lab/.venv/bin/python -m pip
install --upgrade pip

Note: you may need to restart the kernel to use updated packages.

```
[nltk_data] Downloading package punkt to /Users/almadani/nltk_data...
```

```
[nltk_data] Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package punkt_tab to
[nltk_data] /Users/almadani/nltk_data...
```

```

('moses',)
('supposes',)
('his',)
('toeses',)
('are',)
('roses',)
('but',)
('moses',)
('supposes',)
('erroneously',)

```

```
[nltk_data] Unzipping tokenizers/punkt_tab.zip.
```

Calculating probabilities in unigram models

```
### Calculate counts of each word

string = "moses supposes his toeses are roses but moses supposes erroneously"

unique_words = set(word_tokenize(string))

words_frq = []
for uniques in unique_words:
    frq = string.count(uniques)
    words_frq.append((uniques, frq))

words_frq

[('erroneously', 1),
 ('supposes', 2),
 ('are', 1),
 ('roses', 1),
 ('toeses', 1),
 ('his', 1),
 ('moses', 2),
 ('but', 1)]
```

```
### Calculate probabilities:

probab = []

for count in range(len(words_frq)):
    pro = words_frq[count][1]/len(word_tokenize(string))
    probab.append((words_frq[count][0],pro))

print('The probabilities are: ', probab)
```

The probabilities are: [('erroneously', 0.1), ('supposes', 0.2), ('are', 0.1), ('roses', 0.1), ('toeses', 0.1), ('his', 0.1), ('moses', 0.2), ('but', 0.1)]

Bigram

Bigram an n-gram consisting of two items of a sequence

- Two words that are used together to mean something specific

$$P(B | A) = \frac{\text{Count of } (AB)}{\text{Count of } (A)}$$

Method 1

```
import nltk
from nltk.tokenize import word_tokenize

string = "moses supposes his toeses are roses but moses supposes erroneously"

bigrams = nltk.bigrams(word_tokenize(string))
for i in bigrams:
    print(i)
```

```
('moses', 'supposes')
('supposes', 'his')
('his', 'toeses')
('toeses', 'are')
('are', 'roses')
('roses', 'but')
('but', 'moses')
('moses', 'supposes')
('supposes', 'erroneously')
```

Method 2

```
from nltk.util import ngrams
from nltk.tokenize import word_tokenize

string = "moses supposes his toeses are roses but moses supposes erroneously"

bigrams = ngrams(word_tokenize(string), 2)

for i in bigrams:
    print(i)
```

```
('moses', 'supposes')
('supposes', 'his')
('his', 'toeses')
('toeses', 'are')
('are', 'roses')
('roses', 'but')
('but', 'moses')
('moses', 'supposes')
('supposes', 'erroneously')
```

Trigram

Trigram an n-gram consisting of three items of a sequence

- Three words that are used together to mean something specific

Method 1

```
import nltk
from nltk.tokenize import word_tokenize

string = "moses supposes his toeses are roses but moses supposes erroneously"

trigrams = nltk.trigrams(word_tokenize(string))

for i in trigrams:
    print(i)
```

```
('moses', 'supposes', 'his')
('supposes', 'his', 'toeses')
('his', 'toeses', 'are')
('toeses', 'are', 'roses')
('are', 'roses', 'but')
('roses', 'but', 'moses')
('but', 'moses', 'supposes')
('moses', 'supposes', 'erroneously')
```

Method 2

```
from nltk.util import ngrams
from nltk.tokenize import word_tokenize

string = "moses supposes his toeses are roses but moses supposes erroneously"

trigrams = ngrams(word_tokenize(string), 3)

for i in trigrams:
    print(i)
```

```
('moses', 'supposes', 'his')
('supposes', 'his', 'toeses')
('his', 'toeses', 'are')
('toeses', 'are', 'roses')
('are', 'roses', 'but')
('roses', 'but', 'moses')
('but', 'moses', 'supposes')
('moses', 'supposes', 'erroneously')
```

What is padding in NLP?

Padding is a technique used to ensure that all input sequences have the same length. This is necessary because neural networks require inputs that have the same shape and size. However, when we pre-process and use texts as inputs for our model, not all sentences have the same length. In other words, some sentences are longer or shorter than others. We need to have inputs with the same size, and this is where padding comes in.

Padding is a special form of masking where the masked steps are at the start or the end of a sequence. Padding comes from the need to encode sequence data into contiguous batches: in order to make all sequences in a batch fit a given standard length. The padding is added before splitting the sentence.

n-1 padding is added.

```
from nltk.util import ngrams
from nltk.tokenize import word_tokenize

string = "moses supposes his toeses are roses but moses supposes erroneously"

print('----- Padding in Bigrams -----')
list(ngrams(word_tokenize(string), pad_left=True, left_pad_symbol("<s>",
                                     pad_right=True, right_pad_symbol("</s>", n=2))
```

----- Padding in Bigrams -----

```
[('<s>', 'moses'),
 ('moses', 'supposes'),
 ('supposes', 'his'),
 ('his', 'toeses'),
 ('toeses', 'are'),
 ('are', 'roses'),
 ('roses', 'but'),
 ('but', 'moses'),
 ('moses', 'supposes'),
 ('supposes', 'erroneously'),
 ('erroneously', '</s>')]
```

```
print('----- Padding in Trigrams -----')
list(ngrams(word_tokenize(string), pad_left=True, left_pad_symbol("<s>",
                                     pad_right=True, right_pad_symbol("</s>", n=3))
```

----- Padding in Trigrams -----

```
[('<s>', '<s>', 'moses'),
 ('<s>', 'moses', 'supposes'),
 ('moses', 'supposes', 'his'),
 ('supposes', 'his', 'toeses'),
 ('his', 'toeses', 'are'),
 ('toeses', 'are', 'roses'),
 ('are', 'roses', 'but'),
 ('roses', 'but', 'moses'),
 ('but', 'moses', 'supposes'),
 ('moses', 'supposes', 'erroneously'),
 ('supposes', 'erroneously', '</s>'),
 ('erroneously', '</s>', '</s>')]
```

Alternatively, We can get rid of the parameters by using

pad_both_ends method in NLTK

Note the n argument, that tells the function we need padding for bigrams.

```

from nltk.lm.preprocessing import pad_both_ends
from nltk.util import ngrams
from nltk.tokenize import word_tokenize

string = "moses supposes his toeses are roses but moses supposes erroneously"

list(ngrams(pad_both_ends(word_tokenize(string), n=2), n=2))

```

```

[('<s>', 'moses'),
 ('moses', 'supposes'),
 ('supposes', 'his'),
 ('his', 'toeses'),
 ('toeses', 'are'),
 ('are', 'roses'),
 ('roses', 'but'),
 ('but', 'moses'),
 ('moses', 'supposes'),
 ('supposes', 'erroneously'),
 ('erroneously', '</s>')]

```

Training the Maximum Likelihood Model

Maximum likelihood estimation estimates the model parameters such that the probability is maximized.

- Get the requisite n_grams frequency counts from a corpus.
- Normalize them to a 0-1 range.
- Build the vocabulary (of unique words): to create this vocabulary we need to pad our sentences (just like for counting ngrams) and then combine the sentences into one flat stream of words.

To do the above steps, we need to create an Ngram and add the padding then flatten the list. All these steps are requested for pre-processing to train the MLE model. However, we can use `padded_everygram_pipeline()` in NLTK to do all the steps in one call.

The `padded_everygram_pipeline()` will return:

- 1- **ngrams_gen**: Each sentence becomes a list of n-grams with padding.
- 2- **vocab_gen**: A flat stream of all tokens, with start and end symbols included

References:

<https://www.nltk.org/api/nltk.lm.html> (<https://www.nltk.org/api/nltk.lm.html>)

<https://www.kaggle.com/code/alvations/n-gram-language-model-with-nltk>
(<https://www.kaggle.com/code/alvations/n-gram-language-model-with-nltk>)

```

from nltk.lm.preprocessing import padded_everygram_pipeline
from nltk.util import ngrams
from nltk.tokenize import word_tokenize
from nltk.lm import smoothing

string = "moses supposes his toeses are roses but moses supposes erroneously"

#Create Bigram Language Model
train, vocab = padded_everygram_pipeline(2, [word_tokenize(string)])

```

```

from nltk.lm import MLE
lm = MLE(2)
lm.fit(train, vocab) #Train the Maximum Likelihood Model

```

Calculating Probabilities

Here's how you get the score for a word given some preceding context. For example we want to know what is the chance that "b" is preceded by "a".

```
print(lm.score('supposes', ['moses'])) #Calculate Bigram Probabilities for P('moses'|'supposes')
print(lm.score('moses', ['supposes'])) #Calculate Bigram Probabilities for P('supposes'|'moses')
- Returns 0 because unseen
```

```
1.0
0.0
```

Evaluating n_grams models using perplexity

Perplexity is a measurement of how well a probability model (like a language model) predicts a sample.

It's the probability of the test set normalized by the number of words.

A lower perplexity means:

- The model is less "perplexed" by the data.
- The model assigns higher probabilities to the actual words in the test set.
- Therefore, better performance.

A higher perplexity = more uncertainty = worse model.

```
test = [('moses', 'supposes'), ('toeses', 'are')]
print(lm.perplexity(test)) #Perplexity for bigram model
```

```
1.0
```

Task: Generating tweets using n_grams

In this Task, we will build a MLE model with n-gram to generate tweets.

The agenda is:

1. load the data: <https://www.kaggle.com/datasets/adizafar/large-random-tweets-from-pakistan>
(<https://www.kaggle.com/datasets/adizafar/large-random-tweets-from-pakistan>)
2. pre-processing the data:
 - Remove hashtags
 - Remove RT
 - Remove any website
 - Remove any mentions
 - Remove emojis
3. build MLE model with n-gram (Bi-gram)
4. evaluate the model
5. calculate the probability of the bigram: (pakistan is)
6. calculate the perplexity of the word: (pakistan)

before we start we need to install library to handle the emojis in tweets

```
pip install emoji
```

```
import pandas as pd
import re
import html
import emoji
from pathlib import Path
from nltk.lm import MLE
from nltk.lm.preprocessing import padded_everygram_pipeline, pad_both_ends
from nltk.util import ngrams

folder = Path('lab3') if Path('lab3').is_dir() else Path('.')
df = pd.read_csv(folder / 'pakistan_tweets.zip', encoding='utf-8',
                  encoding_errors='replace', low_memory=False)
df[['full_text']].head()
```

	full_text
0	تیرا لیٹر میرا لیٹر نواز شریف نواز شریف کیا
1	Happy birthday to my brother n boss , May you ...
2	
3	`suspicious °jikook au jimin'in yaşadığı kasab...
4	Speaking of @SpiderMan... https://t.co/uGJ...

```
def clean_tweet(text):
    text = html.unescape(text)
    text = re.sub(r'https?://\S+|www\.\S+', ' ', text, flags=re.I)
    text = re.sub(r'#\w+', ' ', text)
    text = re.sub(r'@\w+', ' ', text)
    text = re.sub(r'\bRT\b', ' ', text, flags=re.I)
    text = emoji.replace_emoji(text, replace=' ')
    return ' '.join(re.findall(r'\b\w+\b', text.lower()))

cleaned = df['full_text'].dropna().apply(clean_tweet)
cleaned = cleaned[cleaned.str.len() > 0].drop_duplicates()
cleaned.head()
```

```
0    تیرا لیٹر میرا لیٹر نواز شریف نواز شریف کیا آپ...
1    happy birthday to my brother n boss may you ha...
3    suspicious jikook au jimin in yaşadığı kasabad...
4                                speaking of
5                                alexa skip to 2021
Name: full_text, dtype: str
```

```
tweets = cleaned.sample(frac=1, random_state=42).str.split().tolist()
split = int(len(tweets) * 0.8)
train_tweets = tweets[:split]
test_tweets = tweets[split:]

train, vocab = padded_everygram_pipeline(2, train_tweets)
model = MLE(2)
model.fit(train, vocab)
```

```
test_bigrams = []
for tweet in test_tweets:
    test_bigrams.extend(ngrams(pad_both_ends(tweet, n=2), 2))

model.perplexity(test_bigrams)
```

inf

```
model.score('is', ['pakistan'])
```

0.042269666141280984

```
model.perplexity(['pakistan',])
```

287.72316950852564

```
for seed in [1, 2, 3]:
    generated = model.generate(25, text_seed=['<s>'], random_seed=seed)
    words = []
    for word in generated:
        if word == '</s>':
            break
        if word not in ['<s>', '<UNK>']:
            words.append(word)
    print(' '.join(words))
```

can t sleep all our hearts of the big a stay i still 15 km long day when you

چین کے مابین اشتراک اور پر نواز شریف دی نا ممکن ہو اگر دو سال چرس بیرونی کی صورتحال میں
بہتری آئی

h ghum jati half of all